

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

THIRD YEAR INDIVIDUAL PROJECT

LOST:
The Logic Semantics Tutor

Author:
Alina BOGHU

Supervisor:
Ian HODKINSON

Second marker:
Fariba SADRI

June 16, 2013

Abstract

The aim of this project was to develop a software tool that can teach the semantics of first order predicate logic to students by helping them visualise the process of sentence evaluation. Thus, the focus was on developing an intuitive and engaging user interface to show and allow modification of structures, signatures and sentences, as well as provide relevant exercises for the student to practice with. The latter is arguably the most important feature of this tool and an addition to the functionality of the previous LOST. The user can now ask to solve 3 automatically generated types of puzzles, within a tutorial environment that measures their progress. Completing each lesson is an actual achievement and provides real confirmation of understanding the semantics of first order logic.

For robustness, the tool is linked to a lesson database which can be accessed by students and lecturers alike. The former also have permissions to add, edit or remove lessons and see their student's progress.

I believe these are firm grounds for many possible extensions (such as a Hintikka game) and can be of real use, standalone or alongside the first year predicate logic course. This report will provide further detail of its implementation and purpose.

Acknowledgements

I would like to express gratitude and appreciation to my supervisor, *Ian Hodgkinson*, for his continuous support and motivation that have guided me towards completing this project, as well as to my second marker, *Fariba Sadri*, for her useful remarks and suggestions.

Altogether, the resources, staff and students of the Department of Computing have made working on this project enjoyable every day.

Contents

1	Introduction	3
1.1	Motivation and relevance	3
1.2	A brief, non-technical description	3
1.3	Report Overview	4
2	Background	5
2.1	First order predicate logic semantics	5
2.2	Existing solutions	6
2.2.1	LOST 2007	7
2.2.2	Tarski's World	7
2.3	Ideas shaping the final approach	8
2.4	Software choices	8
2.5	Software Architecture Design Patterns	10
3	Approach and Implementation details	11
3.1	Evaluator	11
3.1.1	Logic Tree	12
3.1.2	Signature	14
3.1.3	Structure	15
3.2	Parser	15
3.3	Adaptor	16
3.4	User Interface	17
3.4.1	The View	17
3.4.2	The Controller	18
3.5	Exception handling and user guidance	18
3.5.1	Instructions, tips and tricks	18
3.5.2	Error messages	18
4	Evaluation	19
4.1	Quantitative evaluation	19
4.2	Qualitative evaluation	19

5	Conclusions and Future Work	20
5.1	Learning outcomes	20
5.2	Potential improvements	20
5.3	Potential extensions	20
A	Code UML Diagram	22
B	Parser grammar	23
C	Short demo	25

Chapter 1

Introduction

LOST stands for LOgic Semantics Tutor and aims at providing students with a helpful tool for learning the semantics of first order predicate logic.

1.1 Motivation and relevance

First order logic is a powerful tool, of great importance in computing, mathematics and really any system relying on proofs. For students it is a way of practising and developing important logical skills, it provides a formal mean for studying and understanding common mathematical structures and it forms a rigorous mind. However it is difficult for most students to learn the semantics of first order logic and with this being the key to assigning a truth value to any first-order logical sentence, the issue is important. This project attempts to provide a simple software solution and a solid base for future extensions.

On the current market, there are a considerable number of tools designed to help teach natural deduction or equivalences, but very few attempts have been made to design something that will accompany the student in learning the semantics of first order logic and even fewer of them are free. Feedback is currently restricted to lectures and tutorials which means this kind of instant guidance provided by computer software would make a great impact in the way students learn, improving both their and the teacher's experience. Designing such a tool raises interesting questions of what would make an interface intuitive and what does it need to provide the user with, such that they can learn and experiment with the semantics of first order predicate logic.

1.2 A brief, non-technical description

After a bit of thought and research into current solutions, I realised it would not be easy to implement in an attractive way and this will become more obvious after discussing the existing work. Thus one of my guidelines shaped up

to be engaging the student as much as possible which software can do if it is intuitive, bug free and most importantly rewarding. The first of these would be achieved if the user is allowed to naturally interact with the application. Because in reality no one likes to read manuals, I aimed to make my program safe to use with a “click and see what happens” approach. This required reasoning about as many occurable events and exceptions as possible. As for the actual content point of view, it would be useful if the user could visualize structures in such a way that they could easily guess the meaning of a possible sentence within it. They should be able to use a toolbox-like signature representation to add and remove objects and relations as well as have a way of introducing and evaluating sentences. At this point it became clear that the interface would have three main aspects to handle: structures, signatures and sentences, which will be discussed further in the implementation details section.

Next, I aimed for a solid back end that could handle even a completely different interface implementation. This meant for it to be able to correctly evaluate the semantics of a logic formula given a structure, regardless of the input or output method. Parsing the input sentence (be it from a terminal or a GUI) would therefore have to be done reliably and its representation in memory be clear, effectively accessible and easy to evaluate.

Features described so far would constitute a minimum viable product, however making it such that the student feels their time was spent worthily was just as important. I aimed to achieve this by implementing a simple quiz system with the possibility of further development. This particular functionality would provide the confirmation a student needs when asked to learn anything.

1.3 Report Overview

In the following chapters the reader can expect a walk through the process of developing the final product, from gathering background information to implementation details, testing procedures and, as much as possible honest evaluation of the outcome. Though previous knowledge of first order predicate logic and Java would help in understanding the effort that went into this project and its relevance, it is not compulsory. I will revise the essential aspects of the theoretical side, however basic knowledge of logic operators and their meaning (i.e. and, or, implies etc.) is assumed.

Chapter 2

Background

2.1 First order predicate logic semantics

In order to understand the product I am aiming for we should first take a look at what first order predicate logic is and why its semantics can be tricky. As an extension of propositional logic, it expresses statements such as *Socrates is a man* in much more detail. While propositional logic would regard this sentence as atomic and simply assign it a truth value, predicate logic provides a way of describing its internal structure and of evaluating it within a relevant context. I will use this sentence to briefly introduce the key concepts used in predicate logic to make “splitting the atom” possible. These are:

- *Constants* - which name the objects inside a context (e.g. Socrates). One constant can name exactly one object.
- *Relation symbols* - which describe properties of the objects they take as arguments or, in the case of nullary relations (which take no arguments), general properties of the structure (e.g. *man* is a unary relation, it represents a property that Socrates may or may not have).

In order to keep track of these two new concepts, we use a *signature*. This represents the syntactic side of evaluating a sentence and provides the necessary tools: constant names and relations symbols. For a computer scientist it may be easier to look at it as a collection of abstract classes that can be instantiated to form the structure and give it meaning.

Using just these concepts, we can now rewrite the sentence we discussed as *man(Socrates)*. However we still cannot decide its truth value and at this point two questions arise: First, which is the object that Socrates describes? Second, what does it mean for something to be a man?

This is where the *structure* comes in. It is defined to be a non-empty set of

objects that the signature knows about. If we take our structure to be an imaginary world of hobbits and name one of them Socrates, our sentence would be false, as Socrates would not be a man, he would be a hobbit. However in the context of the real world where Socrates names the famous philosopher, the sentence is true.

Next, if we want to express something like *All men are mortal* we must introduce the two quantifiers. Again, these can be viewed as a way to iterate over the objects that form a structure. These are:

- \exists (*Exists*) - which checks that there is at least one object in the structure that satisfies the sentence it refers to and makes it true.
- \forall (*For all*) - which checks that all of the objects in the structure satisfy the sentence it refers to and each make it true.

The sentence can now be written as $\forall (men(x) \rightarrow mortal(x))$ and we refer to x as a bound variable because it is pinned down by a quantifier, in this case \forall . A more precise reading of the above formula would be: "If something is a man than it is mortal."

Another aspect that the user must understand is that sentences containing unbound variables are also valid but cannot be evaluated to a truth value. Saying " $men(x) \rightarrow mortal(x)$ " makes no sense until we decide what x refers to. If x were a constant than it would refer to the object it names, however common practice dictates that we reserve the last letters of the alphabet for naming variables. It is this kind of subtlety that I am hoping to make clearer with the help of an interface which shows exactly which objects are named by constants and which can only be referred to by using quantified variables.

The theoretical side is obviously what forms the base of this project. It is therefore natural that the implementation follows its key aspects: the structure, the signature and the logic formulas. Throughout the report I will inevitably make references to the concepts summarised in this section whilst possibly making further additions and clarifications.

2.2 Existing solutions

As mentioned before, previous attempts have been made to provide software solutions to teaching predicate logic semantics. In fact one student provided a solution for this same project specification in 2007. Furthermore, the Openproof Project at Stanford's Center for the Study of Language and Information (CSLI) is concerned precisely with the application of software in logic and they have developed Tarski's World for this purpose.

2.2.1 LOST 2007

This application was previously available on the DoC's lab machines. Currently however, the only available resource is the user manual. As it is a solution to the same project specification I have carefully studied it and picked up what I thought were several good ideas, whilst taking note of things I should avoid.

The application provides a good representation of the logical structure. It displays objects as circles, filled with colours corresponding to the unary relation symbols that apply. Binary relation symbols are represented as arrows, also colour coded. The user can drag objects to rearrange them as he seems fit. This representation is quite clear and easy to interpret. For this reason my own implementation is similar, with a few changes aimed at smoothing the experience even further. First I decided objects should have a clear base colour displaying only their name in the case of constants. Then, according to which unary relations apply to it, coloured borders would be added or removed. This eliminates the risk of a confusing pie chart when an object has many relations referring to it. Next, the arrows would be labelled with a list of names of the relations they represent, such that if multiple relations exist amongst two objects there is no need for multiple arrows rendering better space efficiency.

Furthermore the application allows the user to interact with the structure with several buttons, a signature tree and a number of forms for creating structures and introducing formulas for evaluation. These are offered to the user according to their intention. The general purpose of all this is important and the functionality essential to the application. However in my own implementation I tried to minimise the number of buttons and additional windows or forms, keeping it simple and in one place.

Finally the application offers the possibility to play a Hintikka game, which is a wonderful way of practising logic semantics skills. It also makes an attempt at logic to English translation. These would make very useful extensions to my own application as I chose to focus on the tutorial aspect.

2.2.2 Tarski's World

The Tarski's World application is based on the same principle. Its representation of the structure is a 3D world of blocks. It uses an interpreted first-order language which allows users to write sentences about the world and evaluate their truth. A Henkin-Hintikka game is also provided, along side the main game which consists of questions about the active sentences. This was useful when implementing my own tutorial questions.

My single reservation was once more the number of menus, panels and buttons that can become quite irritating since only a few are truly relevant and most don't seem worth the effort of figuring out. From the point of view of a

student this can be quite off putting and might drive them towards the classic pen and paper approach. Otherwise it is generally very professionally made and I didn't encounter any major flaws. The learning experience is complete if we take into account the extensions including text books and an online evaluation system that grades the student's performance which I particularly appreciated.

2.3 Ideas shaping the final approach

After looking into the above mentioned solutions and others similar to them it became obvious that the main difficulty with providing a teaching tool for logic semantics is in fact teaching the student to use it. When working with these tools I got the feeling they let the back end lead the implementation of the interface instead of using the latter to hide the inner works. For example, just because there are say 20 relation symbols active in the signature doesn't mean the user wants to see them at all times in a big grid of buttons that takes up half of the work space. I believe it is important to let the user focus on semantics without having to worry too much about the syntactic aspect, which should be provided and hidden as much as possible.

Another remark should be made on the tendency of developers to forget the user doesn't know as much as they do. It is impossible to develop a teaching tool without understanding first order predicate logic inside out, which makes it is even harder to empathise with a student that is just starting to learn it. It is therefore easy to unconsciously overestimate the user's knowledge and forget to provide explanations that might in fact be useful to them. For example it might seem counter intuitive that a sentence such as `man(Socrates)` is false. But if the signature's only unary relation symbol is *hobbit* this may seem obvious to a logician. However the user who tried to input the above sentence might benefit from a message such as "The unary relation symbol `man` is not defined in this structure", instead of their sentence just being rejected as invalid.

Finally, I fixed the two features that would make this project unique. First of all I would keep it as simple as possible, making it a priority to minimise the number of buttons and choices the user has to make. This had to be done without sacrificing any of the essential functionality. Secondly, I would make it rewarding, by giving the user an opportunity of testing their skills. As several students have worked on this project specification across the years, I figured this would be a useful and fresh addition.

2.4 Software choices

Having sufficiently stressed the importance of the interface I had to make a choice as to which programming environment would best fit its requirements. I wanted to be able to focus on building something attractive without having to

worry too much about the lower level issues like memory assignments, garbage collection or optimal for loop nesting. This, as well as my previous programming experience, perused me to choose Java and Swing, which proved to be at a high enough level for my interface whilst offering me the freedom to design the back end for the sentence evaluation from scratch. Thus I had complete control and understanding of every step that takes a sentence from it's raw form to an outcome. And understanding made it easier to track bugs, make changes and pursue a steady development. Java also ensures a decent compatibility across computer platforms, has extensive documentation and is a reliable language that addresses a large audience.

Of further help with the the interface was the IDE. After starting off with Eclipse, I soon realised the simple task of arranging buttons on the panel was uselessly tedious. So after a little research I chose to use NetBeans for its GUI Builder. It took a bit of getting used to, especially with learning to get around the read-only generated code, but eventually everything proved to be 100% customisable and certainly worth switching to. Overall this saved me significant time and helped me achieve my desired esthetic standard.

Java proved to be the right choice again when it came to writing a parser for the logic formulas. It made it easy to link an automated parser, reducing my task to writing a comprehensive grammar. At this point, another choice was to favour Antlr4 over the previous, better known version. What convinced me was its ability to deal with left recursive grammars. Although removing recursivity is a straight forward algorithm, it makes the grammar quite ugly and a bit more difficult to read. Antlr4 also uses an impressive adaptive parsing strategy, making it significantly more time efficient than its static ancestors.

For version control I used Git. It is an easy and reliable way of keeping backups, working remotely, tracking progress and proved essential for a project this size as it happened more than once that I had to revert to previous commits and slowly add in changes to discover bugs. It also proved useful for its branches, allowing me to experiment with external Java packages such as mxgraph.

For the report I used LaTeX with Vim, which greatly facilitated the text formatting and made it easy to keep track of the different sections. The Unix shell was also of great use with compiling Antlr and LaTeX and using Git as well as with occasional remote connections. Finally, the entire project was developed on a Ubuntu 12.04 system.

2.5 Software Architecture Design Patterns

As developing proceeded, the size of the code became overwhelming. The separations packages and classes provided was insufficient. When it came to implementing the smallest changes, dependencies inside the code made it a tedious and time consuming task. Therefore, when the back end was finished and I began developing the user interface, I took a step back and revised some software design patterns, in order to choose a system that would best fit my requirements.

It was the *Model-View-Controller* pattern that helped the most. By completely separating the back end from the display I was able to keep better track of the whole picture:

- Model - this represents everything in the back-end. It handles all the information concerning sentences, signatures, structures and evaluation.
- View - this represents the interface. It contains all the component types needed to represent logic objects. It also handles the display layout and appearance and contains the project's main method.
- Controller - this is the class that links the model to the view. To do this it pairs up logic objects with an appropriate graphical component. It also interprets user input such that it can be sent back to the model for interpretation.

Having separated the information from the representation in this way, it became much easier to change the interface since I no longer had to worry about the compatibility with the model. This was essential as designing the GUI meant repeatedly using the application and applying slight variations to see what would be most comfortable for the user.

I also used the *Adaptor* structural design pattern, which allowed me to separate the parser generating from the sentence structure. I then used an the Sentence class constructor to convert the parse tree into my own logic tree, which I had designed with an integrated evaluator.

Chapter 3

Approach and Implementation details


This logic semantics tutoring tool is naturally shaped around the theoretical elements that provide logic formulas with meaning. It contains three packages which I will discuss in the same order they were developed:

- Evaluator - that holds the logic tree structure along with an adaptor for the parse tree and an evaluator.
- Parser - which handles the translation of raw logic formulas input as strings into a parse tree.
- GUI - which contains the user interface and controller.

3.1 Evaluator

As it can be deducted from the brief description above, this is the main and largest part of the software. I will start with how logic formulas are supported then move onto the evaluation process and finally the adaptor.

The starting point of my project was representing sentences in an efficient computer readable format. It was also important to remain close to a human friendly representation. The solution came once from logic theory, where sentences are often represented with the help of logic trees. These trees have logic operators or quantifiers for nodes and simple atomic formulas for leaves. For example, the sentence $\forall (men(x) \rightarrow mortal(x))$ is represented by the following tree:



So with this in mind the first step was to build the classes to represent the logic elements that would be used to form the nodes of the logic tree.

- NullaryRel, UnaryRel, BinaryRel
Classes to represent relations symbols with arities from 0 to 2. Although it is possible to have higher arities they are not common in predicate logic and would over complicate the representation too much. Once a students understands these three, relations of higher arities are merely an inductive step away.
- Term
A class to represent the parameters relation symbols take. It has a name and is extended by Constant and Variable. The latter adds boolean fields indicating whether it is bound by a certain quantifier.
- BinOp
An enum type containing all the binary logic operators along with definitions of their behaviour (i.e. functions overriding the abstract function evaluate).

3.1.1 Logic Tree

Next I build the actual nodes. Each extend the abstract LogicTreeNode class, overriding the evaluate function and providing a pointer to the next node (or nodes in the case of binary operators). They are of course encapsulated by a LogicTree class which has the head of the tree to nicely call evaluate upon. Let us see how each node works in decreasing order of their complexity:

- ForAllNode
This node's evaluation function returns the boolean value representing whether each of the objects in the structures make the sentence true. And it is very important to understand the difference between each object making the sentence true and all objects making the sentence true. For example, to verify the sentence $\forall x (men(x) \rightarrow mortal(x))$ (All men are mortal) we would take each man in turn and verify that he is mortal, disregarding any animals or other things. It would be wrong to first verify

that everything is a man and if so, then verify that everything is mortal. Such a sentence would be written as $\forall x \text{ men}(x) \rightarrow \forall x \text{ mortal}(x)$. To make this even more clear, let us assume the structure contains an immortal man. This would make the first sentence false because not all men are mortal. However in the case of the second sentence, when verifying that everything is a man, we would find that is not true since there are also animals and other things in the structure. And since falsity implying anything is always true, the second sentence is true. Hence the two sentences have different outcomes and mean different things.

In order to handle this correctly, the forall node passes an assignment down the tree to be evaluated with the entire sentence in scope. An assignment represents an object from the structure's domain and to build it, the node has a Term field which points to the variable it quantifies. It uses this term to call the Assignment class constructor, then adds the new assignment to an array of assignments as other quantifiers may be encountered along the way. The node passes the assignment down as a parameter to its evaluation method which iterates over all of the objects in turn and only then puts the outcomes together. If all of them verify the sentence then the final result is returned as true, otherwise as false.

- ExistsNode

This node follows the same principle. The difference is outcomes of each assignment are put together with an or operator, such that as soon as one is found to be true, the evaluation stops and returns true. If however it reaches the end of the structure domain without finding a term to satisfy the sentence, it returns false.

- BinaryRelNode

This is a leaf node class and contains a BinaryRel field that holds all the information about the relation that needs to be interpreted. At this point it must be clarified that there are in fact two abstract evaluation functions as follows:

```
abstract boolean evaluate(Signature s, Assignment a);
abstract boolean evaluate(Signature s, Assignment a1, Assignment a2);
```

This is because an assignment must be duplicated in the case of a binary relation if both its parameters refer to the same variable (e.g $\forall x \text{ loves}(x,x)$), or, if both arguments are different variables, naturally an assignment is needed for each.

In the case of constants things are more straight forward: the method simply checks the sentence for the object named by the constant.

- UnaryRelNode

This node works in the same way, with only one assignment for the relation's single parameter.

- **NullaryRelNode**
This node has a `NullaryRel` field and always return's that relation's boolean value.
- **EqualsNode**
Although it may seem like it could be part of the `BinOpNode`, this node is actually a leaf as the equals operator takes terms as arguments. So really it is more similar to the `BinOpNode`. It also requires two assignments in case both arguments are variables and will return whether the two represent the same object.
- **BinOpNode**
This is an internal node containing one of the types of operators defined in the enum type described before, as well as pointers to the nodes it takes as arguments. Its evaluation function calls the evaluation functions of these two nodes in the way described by the operator definition. For example AND will make sure both evaluation functions return true.
- **NotNode**
Also an internal node that points to the rest of the sentence through a `LogicTreeNode` field and simply returns its negated outcome.
- **TruthNode and FalsityNode**
These nodes always return true and false respectively. They represent leaf nodes.
- **DummyNode**
This node is used by the adaptor to translate parse trees and will be further discussed in the relevant section.

3.1.2 Signature

In order to understand the symbols used in the logic formula, a signature is needed. In a series of `ArrayList` fields it holds `Strings` representing names of the `Constant`, `NullaryRel`, `UnaryRel` and `BinaryRel` instances. In order to fill these in, its constructor takes a signature object as a parameter.

The choice for an array list (here and in other classes) was based on the fact that, being based on a dynamically resizable array, it would greatly facilitate adding and deletion of objects as well as locating and iterating over its elements.

Using just the relation symbols' and constants' names it provides another useful check of the formula's syntax before passing it on to the evaluator. If any symbol used in the sentence is missing from the signature and is not a quantified variable, the user will be prompted with the appropriate message. However this will be discussed further in the section regarding interface implementation where the signature's contribution is greater. In the evaluation context its role is simple: it checks that all the symbols used are declared and valid.

3.1.3 Structure

Now that we have seen the representation of a logic formula and signature, we can proceed to understanding the format of the context in which it will be interpreted. Objects constructed by the Structure class are passed as a parameter to the evaluation method such that an interpretation can be made.

The fields of this class are array lists of terms and relations. In order to fill them in, the constructor calls a generate method. This in turn uses four enum types with names for constants and the three types of relations. Each enum type has a random chooser method to return one of its members. The generate method uses this and ensures no duplicates are allowed into the final content of the structure. It generates a random number of actual terms and relations using the Term, Const, and relation class constructors described before.

Every time the program starts up it creates a new structure which the user can modify, thus ensuring a non empty set of elements at all times to remain consistent with the theoretical definition. The details of this will be discussed in the user interface section.

3.2 Parser

I have already discussed the choice of Antlr4 for the purpose of parsing a logic formula from plain string input. In order to generate a parse tree, Antlr needs two things: a grammar and lexer rules.

The grammar has proved quite tricky to write as I wanted to offer the user as much input freedom as possible. I also managed to preserve the operator precedence such that it would make it easier for the adaptor to interpret the parse tree. The final and best solution was once more the simple one, with the main rule shaping to be comprehensive and easy to read:

```
formula
: TRUTH
| FALSITY
| term EQUALS term
| relation
| quantifier formula
| NOT formula
| formula AND formula
| formula OR formula
| formula IMPLIES formula
| formula EQUIV formula
| LPAREN formula RPAREN
```

Please refer to the annex for the complete grammar file.

The **lexer** rules implied making a few decisions as to what the user should be allowed to use for naming the signature elements. For variable names, the general convention is to use letters towards the end of the alphabet, optionally followed by the character `s` and/or a number (e.g. `x3`). I stuck to this convention, however allowing all letters of the alphabet. The rule is as follows:

```
VARIABLE: [a-z] 's'? [0-9]* ;
```

For naming relation symbols I allowed any combination of characters beginning with a lower case as long as its length is greater than 1 or beginning with an upper case letter in which case its length must only be greater than 0:

```
NAME: [A-Z] | [a-zA-Z] [a-zA-Z09'_' ]+ ;
```

The lexer ignores white spaces or new line characters:

```
WS: [ \t\r\n]+ -> skip ;
```

Finally, it provides rules for the logic symbol tokens:


```
LPAREN    : '(' ;
RPAREN    : ')' ;
AND        : '' ;
OR         : '' ;
NOT        : '' ;
IMPLIES    : '' ;
EXISTS     : '' ;
FORALL     : '' ;
TRUTH      : '' ;
FALSITY    : '' ;
EQUALS     : '=' ;
EQUIV      : '' ;
```

Together, the rules above represent the entire lexer file.

3.3 Adaptor

At this point we can return to look at the `LogicTree` class constructor which works as an adaptor for the parse tree. Antlr provides a user basic interface which can be run from the terminal to visualise this tree. For the sentence *All men are mortal* the following tree will be generated: Each node is created from one of the parser context classes.

```
folParser$ConditionContext
folParser$FormulaContext
folParser$QuantifierContext
folParser$RelationContext
folParser$BinaryargContext
folParser$UnaryArgContext
folParser$TermContext
```



The first thing to notice is that the parser will generate a *formula* node that does not correspond a type of logic node but is just a rule that guides the parsing. For this reason I created the dummy node which simply points to the rest of the tree through a *LogicTreeNode* field and returns that node's outcome without altering it.

The adaptor provides a rule of interpretation for each of these classes through a switch statements. The default case is an error message for a complete picture, although it is not reachable, as each node of the parse must be an instance of one of these classes. The constructor performs a breadth first search and generates the appropriate *LogicTreeNode* for each node in the parse tree. Once finished it returns a pointer to the head of the tree. The sentence is now complete and ready for evaluation.

3.4 User Interface

3.4.1 The View

The layout of the workbench is also designed around the three main elements involved in semantics:

- Structure panel
It is the main working area, taking up most of the upper left side. It displays objects as squares with labels for Constants and colour coded

borders to represent unary relations that apply. Binary relations are represented by arrows between objects, also labelled with the list of relations that apply to that arrow. Nullary relations are displayed as a list of labels and can be rearranged to make best use of the space available.

- **Signature panel** It is a tabbed panel at the right side of the workbench, with a tab for each type of element: Constants, nullary, unary and binary relation symbols. Each tab also contains buttons for manipulating their content.
- **Sentence panel** Situated at the bottom of the window, it provides the means for introducing logic formulas and a scrollable list to store them and their outcome.

3.4.2 The Controller

3.5 Exception handling and user guidance

Even the most intuitive interface needs to be able to guide its user if necessary. This is done with the help of useful error messages and clear, concise instructions.

3.5.1 Instructions, tips and tricks

At the top of the window there is of course a menu bar which provides further assistance for the user.

Firstly, it can be used for saving and loading workbenches and their elements. Secondly, it provides a number of tips and tricks to help the user get familiar with the interface.

3.5.2 Error messages

In order to handle bad user input or illegal operations I designed a number of exception classes.

- `DuplicateDefinitionException`
- `UnboundException` and `ThisUnboundException`
- `UndefinedRelationException`

Further error messages are also displayed for inbuilt exceptions:

- **Index out of bounds**
This can occur when a user presses buttons that require a list element to be selected so that the method can pass down an index. In such cases an

error message is displayed only once to instruct the user that they must make a selection.

- I/O exception
This can occur when the user tries to load a corrupted or otherwise inaccessible file. The error message informs accordingly.

Chapter 4

Evaluation

4.1 Quantitative evaluation

```
x(R(Fred,x) y(R(x,y) P(y)))
xy (A B) B D E F G H
xy (A B) B D E F G H
xy (A B) B (D) E (F G) H
A B C
B (D) E H E
Loves(fred,wilma) is_in_the_air(love)

brian = rian
xys(happy(x) A)
x(Red(Fred,x) E E E H E y(R(x,y) P(y)))
A B C D E (Assumes A B C means (A B) C)
A B C (D E)
P (Q R) S
x P(x) y Q(x,y)
```

4.2 Qualitative evaluation

Chapter 5

Conclusions and Future Work

- 5.1 Learning outcomes
- 5.2 Potential improvements
- 5.3 Potential extensions

Bibliography

Appendix A

Code UML Diagram

Appendix B

Parser grammar

```
grammar fol ;

condition
  : formula EOF ;

formula
  : TRUTH
  | FALSITY
  | term EQUALS term
  | relation
  | quantifier formula
  | NOT formula
  | formula AND formula
  | formula OR formula
  | formula IMPLIES formula
  | formula EQUIV formula
  | LPAREN formula RPAREN ;

term
  : VARIABLE
  | NAME ;

quantifier
  : FORALL VARIABLE
  | EXISTS VARIABLE ;

relation
  : NAME binaryarg
  | NAME unaryarg
  | NAME ;
```

```
binaryarg
  : LPAREN term ',' term RPAREN ;

unaryarg
  : LPAREN term RPAREN ;
```

Appendix C

Short demo